

# Homework#2

## CS576 Machine Learning, Fall 2023

**Due: October 12**

### Instruction

- Compile your work into a single file, naming it “*YourLastName\_FirstName\_CS576\_HW2.zip*”
- Structure your submission with individual directories for each part, namely Part I, Part II, Part III, etc.
- For each part, you are required to submit:
  - (1) **Program Code:** Annotate the relevant lines in your code with the corresponding task numbers for clarity.
  - (2) **Execution Results:** Present results for specified tasks, labeling each with its corresponding task number for easy identification.

### Part I. Model Evaluation with Nearest Neighbor Classifier

In Part I, you will conduct a machine learning experiment with **k-Nearest Neighbor** classifier and assess various model performance metrics.

**0.** Initially, download the provided *toydata.csv* provided. This toy dataset incorporates two numeric features (‘Feature1’, ‘Feature2’) and one binary class feature (‘Label’), represented by 0 and 1.

Implement a program which performs the following tasks (a) – (m).

#### 1. [Data Exploration]

- (a) Load the given dataset into your program.
- (b) Visualize the dataset using a **scatter plot**, distinguishing data points by color (e.g., ‘blue’, ‘green’) according to their class label.
- (c) Exhibit the first few rows of the dataset
- (d) Present basic statistics of the descriptive features within the dataset.

#### 2. [Data Preparation]

- (e) Apply the **holdout** method to split the data, allocating **20%** for testing and the remaining **80%** for training. Ensure consistent splitting every time the code executes, by using, for example, “42” as the random seed value in Python.
- (f) **Standardize** the descriptive features in both the training and test sets, excluding the label feature.

#### 3. [Classification]

- (g) Employ the **k-Nearest Neighbor (k=1)** classifier to predict the class label of test instances. Notice that the k-NN classifier does not have a training stage.
- (h) Show all test instances alongside their predicted classes, e.g.,

Test Instances and Predicted Classes:

Instance: [-1.53760567 -2.1124287], True Label: 0, Predicted Label: 0

Instance: [0.93499408 1.15684075], True Label: 1, Predicted Label: 1

#### 4. [Model Evaluation]

- (i) Create a **confusion matrix** for the test data predictions.
- (j) Display **TP, TN, FP and FN** values like the following:

The True Positive (TP) value is: xx

The True Negative (TN) value is: xx

...

- (k) Calculate various model performance metrics such as **Accuracy, Misclassification Rate, Precision, Recall and F1 score** utilizing the TP, TN, FP and FN values. Display the performance measure values like the following:

Accuracy =  $(TP + TN) / (TP + TN + FP + FN) = xx.xx\%$

Misclassification Rate =  $(FP + FN) / (TP + TN + FP + FN) = xx.xx\%$

...

For performance measure calculations, **do NOT use library functions** such as `precision recall fscore support and accuracy scor` in Python. Instead, compute with `confusion matrix values`.

#### 5. [Visualization]

- (l) Generate a **Voronoi** diagram for the training data
- (m) Place the test data points, annotated with their class label, on the Voronoi diagram.

#### Submit the followings:

- (1) A program code that addresses tasks (a) through (m). Please annotate the relevant code lines with their corresponding task numbers for clarity.
- (2) The execution results for tasks (b), (c), (d), (h), (i), (j), (k), (l), and (m). Ensure each result is labeled with its respective task number.

## Part II. Model Comparison with Random Forests

In Part II, you will conduct a machine learning experiment to compare the performance of **Random Forests** under varying hyper-parameter values.

**0.** To start, download the provide **adult.data**. This census dataset originates from a Machine Learning Repository. For a detailed description of the data, refer to <https://archive.ics.uci.edu/dataset/2/adult>. The dataset has a class feature, INCOME, with two categorical values: '<= 50K' and '> 50K'. The task involves classifying whether income exceeds \$50K/yr.

Implement a program to accomplish the following tasks (a) – (l).

### 1. [Data Exploration]

- (a) Load the given dataset.
- (b) Exhibit the initial few rows of the dataset. Show the count of instances and descriptive features in the original data.
- (c) The Adult dataset represents missing values with '?'. Show the count of missing values per each feature.
- (d) Eliminate instances containing missing values. Subsequently, display the updated instance count.
- (e) Illustrate a histogram representing instance counts per INCOME class.

### 2. [Data Preparation]

- (f) The class feature, INCOME, comprises two categorical values: '<=50K' and '>50K'. Transform this feature into binary 0/1
- (g) Implement **One-hot Encoding** for the categorical variables.
- (h) Present the first few rows of the processed data. How many descriptive features does the data now include?
- (i) Allocate **70%** of the data for training, and the remaining **30%** for testing. While splitting the data, ensure that the distribution of classes in the target feature is consistent in both the training and test sets using **stratified sampling**.

### 3. [Random Forests and Performance Evaluation]

- (j) Construct four **Random Forests** models, each varying by the number of trees in the forest (the hyper-parameter m), each **5, 10, 50 and 500**.
- (k) Generate and plot the **ROC curves** of each model collectively.
- (l) Explain which random forests classifier(s) exhibit(s) superior performance.

**Submit the following:**

- (1) A program code that addresses tasks (a) through (l).
- (2) The execution results for tasks (b), (c), (d), (e), (h), and (k), and answer for (l).

### Part III. Hyper- Parameter Tuning with Gradient Boosting

In Part III, you will conduct a hyper-parameter tuning experiment with **Gradient Boosting**.

0. For this experiment, utilize the same adult dataset (**adult.data**) as in Part II.

#### 1. [Data Preparation]

- (a) Load the designated dataset.
- (b) Exhibit the first few rows of the dataset and show the count of instances and descriptive features in the original data.
- (c) Eliminate instances containing missing values.
- (d) The class feature, INCOME, has two categorical values: ‘<=50K’ and ‘>50K’. Alter the target feature to binary 0/1, although it’s generally not a requisite for the Gradient Boosting algorithm.
- (e) Execute **Label Encoding** for categorical variables.
- (f) Illustrate the first few rows of the modified data. How many descriptive features does the data contain? Explain the difference from the prior **one-hot encoding**
- (g) Split the data for model training and testing, allocating **30%** for testing and the remaining **70%** for training.

#### 2. [Hyper-parameter Tuning]

In this experiment, we are primarily altering two hyper-parameters: the number of base learners and the learning rate, for the Gradient Boosting classifier. For **the number of individual decision trees** for the base learners, employ **5, 10, and 50**, and for the **learning rate**, select **0.01, 0.05, and 0.1**. Therefore, a total of 9 combinations will be considered to identify the optimum hyper-parameters.

- (h) Execute a **grid search** to find the most considerable hyper-parameter values among the provided combinations of values. **During the search**, utilize the prepared training data and a **3-fold cross-validation** schema for training and validation. **For testing**, employ the prepared test data, and use **accuracy** as the scoring metric.
- (i) For every combination of the stated parameter values, present the **average test score, standard deviation of test scores, and rank test score** (1, 2, 3..)
- (j) Present the performance report of the model with the superior parameter setting, incorporating metrics such as **accuracy, precision, recall, F1-score**, etc.

**Submit the following:**

- (1) A program code which performs for all the tasks (a) through (j).
- (2) The execution results for tasks (b), (f), (i) and (j), and answer for (f).

## Part IV. Prediction with Linear Regression

In Part IV, you will an experiment with linear regression.

0. Download the provided **auto.csv** file. This dataset will be used to predict the target feature, **mpg**.

### 1. [Data Exploration]

- (a) Load the given dataset.
- (b) Display the initial rows of the dataset and the basic information of features.
- (c) Convert the **horsepower** feature to numeric, if it is not in numeric type within your program.
- (d) Eliminate any features that are non-numeric.
- (e) Utilize **plots** to visualize the pairwise relationships of each descriptive feature with the target feature, **mpg**.

### 2. [Data Preparation]

For this experiment, you will conduct a **simple linear regression** focusing on the **horsepower** feature and the target, **mpg**.

- (f) If any missing values are detected in the 'horsepower', impute them with the feature's **mean** value. For validation, display both original missing values alongside their imputed replacements.
- (g) Divide the data into training and test subsets using the **holdout** method; allocate **30%** of testing, and the remaining **70%** for training.

### 3. [Linear Regression Model]

- (h) Develop a **regression model** utilizing the training data and assess it using the test data.
- (i) Show the **linear regression equation** in the form  $Y=\alpha+\beta X$ .
- (j) Illustrate **the regression line** and data points via a scatter plot. If the intercept is not vividly represented, annotate the intercept value on the plot.

### 3. [Model Evaluation]

- (k) Present a **prediction-true plot** (y-y plot) where the x-axis represent the true target value and the y-axis denoted the predicted target value.
- (l) Evaluate the model using performance metrics: Mean Absolute Error (**MAE**), Mean Squared Error (**MSE**), Root Mean Squared Error (**RMSE**), and Coefficient of Determination ( $R^2$ ).
- (m) Lastly, exhibit a **3D surface plot depicting the error space**, with the x-axis representing the coefficient of 'Horsepower', the y-axis representing the intercept, and the z-axis indicating the mean square error.

**Submit the following:**

- (1) A program code that addresses tasks (a) through (m).
- (2) The execution results for tasks (b), (e), (f), (i), (j), (k), (l), and (m).

## Part V. Support Vector Machine Learning

In Part V, you will experiment with **Support Vector Machine (SVM)** learning.

**0.** For this experiment, you will be utilizing the **toydata.csv** used in Part I. This toy dataset contains two numeric features ('Feature1', 'Feature2') and one binary class feature ('Label'), denoted by 0 and 1.

Implement a program to execute the tasks (a) – (f) as mentioned below.

### 1. [Data Preparation]

- (a) Load the dataset
- (b) Partition the data using the **holdout** method, allocating **20%** of testing, and the remaining **80%** for training.

### 2. [SVM Models]

- (c) Construct four **Support Vector Machine (SVM) models**, each employing different kernel types: **linear**, **poly**, **rbf**, and **sigmoid**.
- (d) Assess the models using the test data.

### 3. [Model Evaluation and Visualization]

- (e) Calculate the **accuracy** of each SVM model
- (f) Per each SVM model, generate a **scatter plot** to visualize the data points in the feature space, coloring them by their true labels. Integrate a contour plot to visualize the **decision boundary** and **margins** of each model.

**Submit the following:**

- (1)** A program code that addresses tasks (a) through (f).
- (2)** The execution results for tasks (e) and (f).